



A.I. LITERACY FUNDAMENTALS

TREVOR YU & CARTER DEMARS



WHY LEARN ABOUT **AI LITERACY**?

Goal: Equip everyone with the tools to understand, talk about, and evaluate AI systems

Why?

- Be able to evaluate and effectively use AI products
- Discern facts from “AI hype” in news and social media
- First steps towards working on AI projects



BY THE END OF THIS WORKSHOP, **YOU SHOULD BE ABLE TO...**

- Classify scenarios as descriptions of artificial intelligence (AI), machine learning (ML) and data science (DS)
- Explain the difference between AI and ML
- Give examples of AI applications
- Explain why data needs to be processed before being used by ML models
- Suggest appropriate types of models for different problems



WHAT IS ARTIFICIAL INTELLIGENCE?

Artificial **intelligence** is machines demonstrating “**intelligent**” capabilities once thought only to be displayed by humans.


- **Intelligence** includes learning, perception, recognition, decision making, conversation, language understanding, planning, control, knowledge representation, reasoning, etc.
- Today's machines are digital computers
- AI include machine **learning**, computer **vision**, automated **decision making**, robotic **control**, **chatbots**, etc.



WHAT IS MACHINE LEARNING?

Machine learning is a computer system that performs a task from data instead of being explicitly programmed.

- ML algorithms are based on the math from linear algebra, calculus and statistics



EXAMPLE: **RECOMMENDATION SYSTEM**

Explicitly programmed (not ML):

- A user selects a short list of tags they're interested in
- If/then/else rules to filter content based on tags
- The user has to explicitly add tags if their interests change

Learns from data (is ML):

- A user "likes" content from their feed, labelling data points of what they prefer
- An algorithm uses the data to learn what things the user wants to see and what they ignore
- The same algorithm can be applied to different users without programming in their individual preferences



SOME COMMON ML TERMINOLOGY

Machine learning is a computer system that performs a task from data instead of being explicitly programmed.

- **Supervised learning**: The system learns a rule to process inputs based on labelled outputs
- **Classification**: Task to separate input examples into two or more discrete groups
- **Regression**: Task to estimate a continuous output value from input data



WHAT IS DATA SCIENCE?

Data science is the use of methods in statistics, computing, and science to extract knowledge from data.

- **Data:** Raw observations about the world
- **Knowledge:** Useful information about the world
- ML techniques are often used for data processing in the data science field



ACTIVITY: **DISTINGUISH BETWEEN AI, ML, AND DS**



MACHINE LEARNING IS A **SUBSET OF AI**

“Machine learning” is often used interchangeably with AI. But:

- AI also includes other techniques like optimization, logic, and knowledge representation
- Machine learning can be studied outside the context of building intelligent systems
- Beyond today’s session, assume we will be using ML techniques to build AI systems



AI, ML, AND DS ARE **CONNECTED**

Human intelligence is natural data processing

- We take inputs like sensory information, memories, and goals
- We output movement, speech, and decisions

Processing and understanding data is fundamental to building intelligent systems

- A data science mindset helps us understand how to build good AI and ML systems
- AI and ML techniques help data scientists draw better insights from data



MACHINE LEARNING IN RESEARCH VERSUS IN PRODUCTION

	Research	Production
Requirements	SOTA Performance	Depend on stakeholders
Computational Priority	High throughput	Fast inference
Data	Static	Constantly shifting
Fairness	Generally, not a focus	Must be considered
Interpretability	Generally, not a focus	Must be considered

Adapted from Designing Machine Learning Systems, O'Reilly.



QUESTION BREAK

What can we clarify about the content so far?



THREE ASPECTS OF **AI LITERACY**

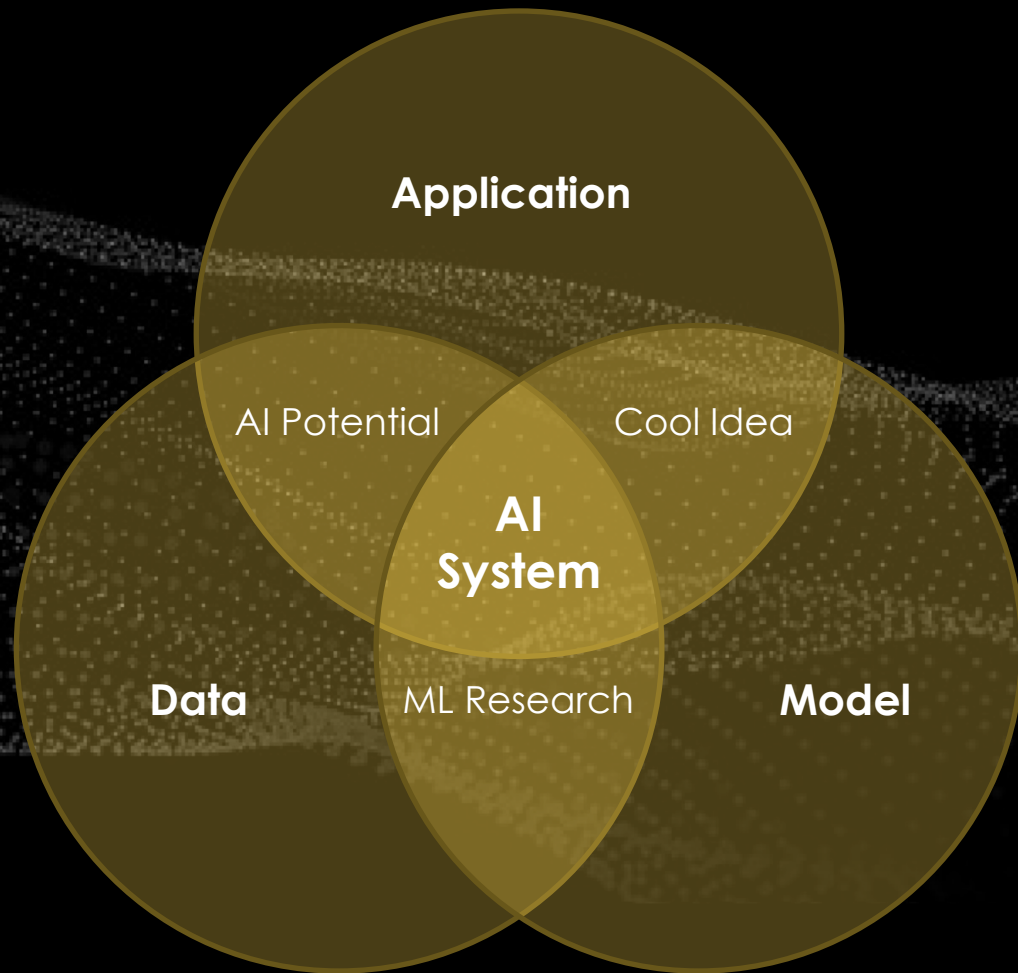
- **Data**: What data is being processed
- **Application**: What intelligent capability is being done and for what purpose
- **Model**: What method is being used to represent intelligence



EXAMPLES OF PAST WAT.AI PROJECTS

- The medical imaging processing project used a **convolutional neural network** to **detect prostate cancer** based on **MRI images**.
- The Open Climate Fix project used a **decision tree** to **estimate solar panel output** based on **weather data**.





UNDERSTANDING DATA

- **Data** is a collection of information about the world, such as facts, observations, measurements, and experiments
- Raw data alone is useless, we need to apply some processing to generate useful inferences about the world




DATA FOR MACHINE LEARNING

- **High quality ML systems require training with a large volume of data**
 - ML models need to learn patterns from many examples
 - Supervised learning also requires labels, which may be hard to collect
 - ML models do better with high quality data
- **Data preprocessing is important because AI models require data to be in specific formats**
 - ML models are based on linear algebra → data should be represented as vectors
 - It takes a lot of time and effort to prepare data for ML



IMAGE PIXELS AS INPUT VECTORS



DATA MODALITIES

- **Images**
 - Natural images, satellite imagery, medical images
- **Tabular**
 - Spreadsheets
- **Text**
 - Books, internet, social media, foreign language
- **Audio**
 - Speech, sonography
- **Time-series**
 - Sensor data, sequences of game actions, video



AI APPLICATIONS: SOLVING PROBLEMS

- **What problem is the AI system trying to solve?**
 - How is success defined and how do we measure it?
 - Is it necessary to use AI to solve this problem?
 - Can the problem be formulated as classification or regression?



AI APPLICATION MODALITIES

- **Tabular data**
 - Identify anomalies, forecast sales
- **Computer vision**
 - Object recognition, image classification, pose estimation
- **Natural language processing**
 - Text classification, entity extraction, language understanding, translation
- **Time-series forecasting**
 - Real-time sensor classification, robotic trajectory estimation
- **Generative**
 - Chat bots, image generation
- **Many industries**
 - Robotics, healthcare, medical science, finance, business, manufacturing, environmental science, education, etc.



LIMITATIONS OF AI SYSTEMS

- We don't yet have artificial general intelligence that can do everything as well as people can
 - Most AI systems are developed to perform a specific task very well
- AI systems still make mistakes
 - AI practitioners need to investigate how and why mistakes occur to prevent and mitigate them
 - Mistakes influence levels of trust in using AI systems
- All AI systems have bias
 - Training data can have bias towards what kind of labels are associated with certain inputs, or there can be gaps in data
 - Assumptions made about the modelling technique can introduce systematic bias towards the way decisions are made
- Large and complex AI systems are expensive to develop and run



AI APPLICATIONS: INTERACTIONS WITH THE WORLD

- **How will the AI system interact with the world?**
 - Consider people, other technology, institutions, laws, conventions, etc.
 - An AI system can perform really well, but if society doesn't use it, it's worthless!



INAPPROPRIATE **USES OF AI**

- **Things that cause harm to people**
 - Warfare, scammer bots, hate speech, discriminatory decision making, privacy invasive surveillance, social manipulation, etc.
- **Situations where there is NO human in the system or when AI would completely replace a human's job**
 - Unless the people being affected themselves want to be replaced
- **When the AI system's performance is poor or inadequately tested**
 - Influences trust in adopting AI systems



ACTIVITY: SUGGEST AI APPLICATIONS

Suggest an application of an existing AI product, or an application of AI that you are excited about



UNDERSTANDING **MODELS**

- A **model** is a conceptual representation of a system and allow us to represent the system in a precise way, such as using a diagram or mathematical notation
 - E.g. modeling the trajectory of a ball using kinematic equations
- **Models make simplifications and assumptions about the systems they represent**
 - Models are never completely correct about their representation, but are often good enough to be useful

An AI model is a simplified representation of “intelligence”. It is not precisely correct about the true nature of intelligence, but it is useful enough to solve problems.



EXAMPLES OF ML MODELS

- **Classic ML and statistical methods**
 - Linear regression, support vector machines, decision trees, probabilistic models
- **Neural networks**
 - Multi-layer perceptron, convolutional neural networks, recurrent networks, transformers
- **Large language models**
 - GPT-4, PaLM, LLaMA



WHEN TO USE WHICH MODELS

Model Type	Data Modality	# Parameters (model size)	# Examples (dataset size)	Tasks
Classical ML	Tabular, time-series, image	10 – 10,000	10 – 10,000	Classification, regression
Small neural networks	Image, time-series, tabular	100 – 100,000	1,000 – 100,000	Classification, regression
Large neural networks	Image, text	100k – 100M	100k – 1B	Classification, regression, low-quality generation
Large language models	Text, image	1B+	1B – 100T	High-quality generation, classification



ACTIVITY: SUGGEST TYPES OF MODELS



ACTIVITY: CHATGPT CASE STUDY

- With a group or on your own, analyze the ChatGPT AI system based on the prompts in the shared document. Write your answers in the shared document for everyone to see.
- bit.ly/3Lo9qml



BY THE END OF THIS WORKSHOP, **YOU SHOULD BE ABLE TO...**

- Classify scenarios as descriptions of artificial intelligence (AI), machine learning (ML) and data science (DS)
- Explain the difference between AI and ML
- Give examples of AI applications
- Explain why data needs to be processed before being used by ML models
- Suggest appropriate types of models for different problems



FUTURE **AI LITERACY TOPICS**

- Designing AI systems for human-AI interaction
- Bias in AI systems
- Reading AI literature



UPCOMING EDUCATION SESSIONS

- Data Preprocessing
- Classical Machine Learning
- Neural Networks for Novices
- Dive into Deep Learning



EXIT SURVEY – ATTENDANCE!

